



AKS Security Best Practices

Alessandro Vozza

Principal Software Engineer @Microsoft, CNCF Ambassador,
Cloud Pirate Captain

Content

- ## Azure Kubernetes Service 101
- ## Security best practices
- ## Recent CVEs
- ## OSS security
- ## Q&A // Discussion

The Pirates

Who are the Cloud Pirates





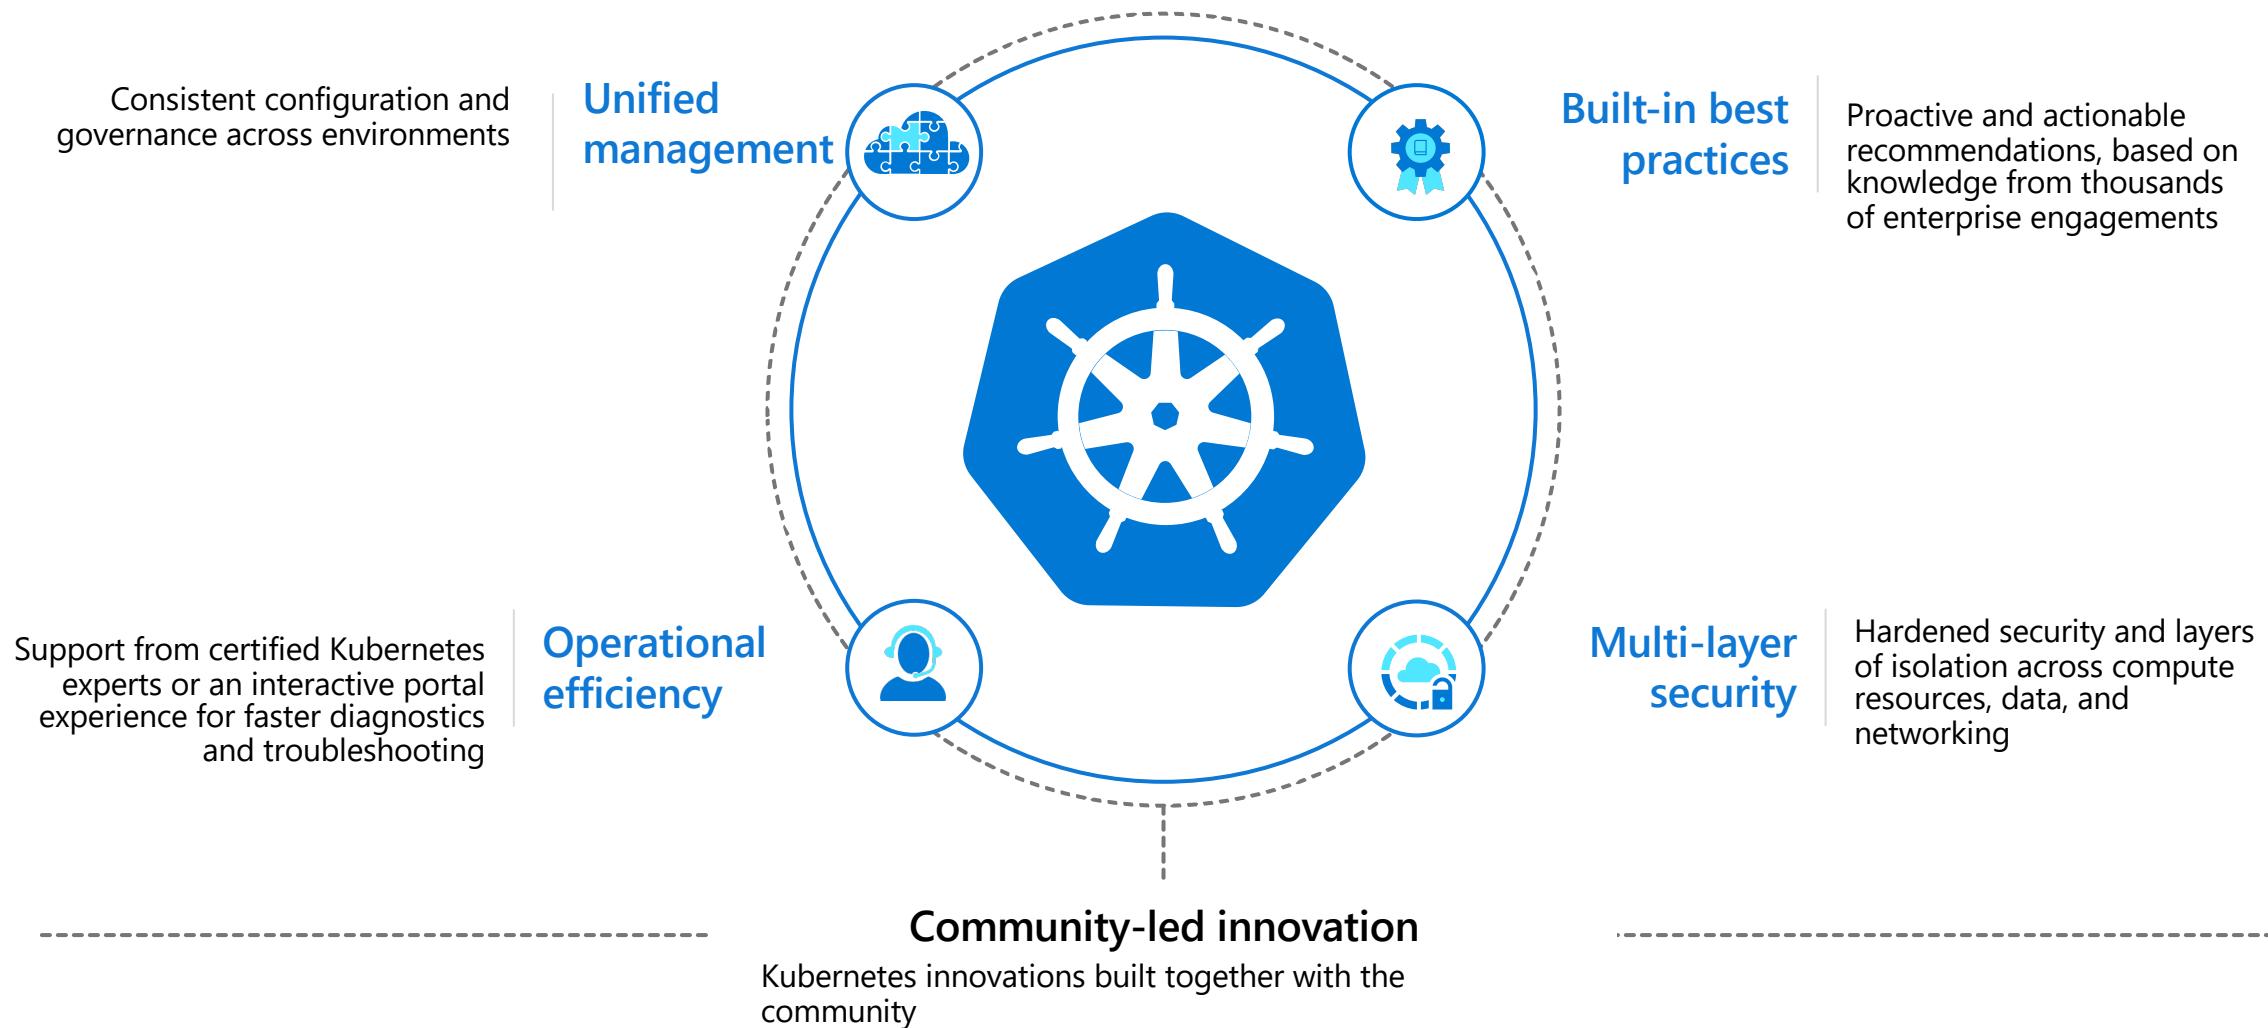
Meet the pirates

The image shows the interior of a modern office space during construction or renovation. The floor is covered in light-colored carpeting. On the left, there's a set of double doors with a red fire extinguisher placed in front of them. A white dry-erase board stands near the entrance. Large windows on the right provide a view of the exterior and some greenery. The ceiling is made of white acoustic tiles, and various pipes and cables are visible on the walls and ceiling, indicating ongoing work.

The Pirate Cove
cloudpirates.nl/signup
Opening mid-October
@cloudpiratesnl

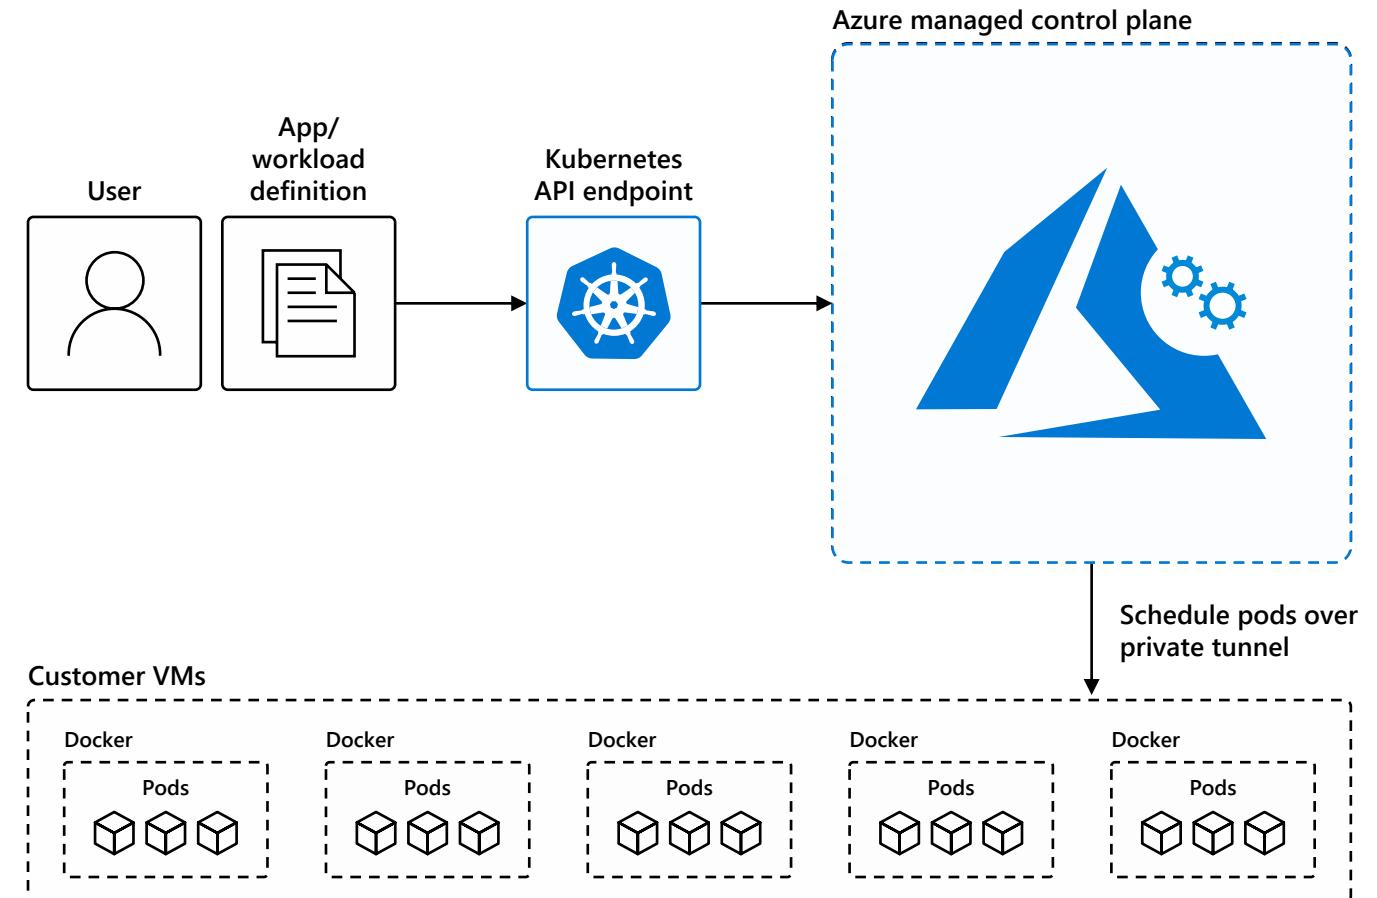
Introduction

Kubernetes on Azure investment areas



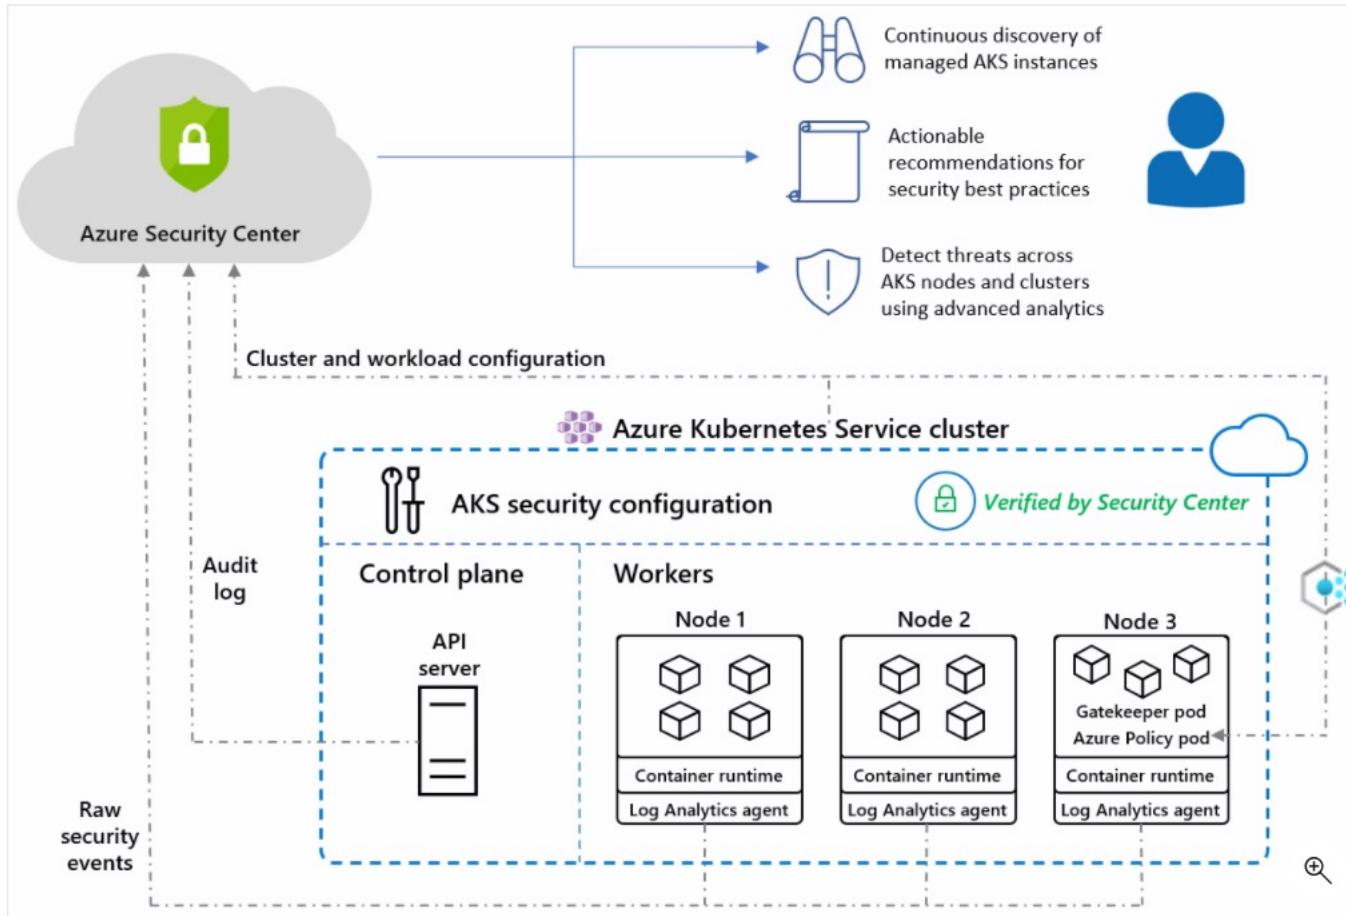
AKS 101

- Automated upgrades, patches
- High reliability, availability
- Easy, secure cluster scaling
- Self-healing
- API server monitoring



Azure Security

Azure Security Center



The screenshot shows the **Security Center | Inventory (Preview)** interface. It displays a summary of resources and filters at the top, followed by three categories of resources:

- Total Resources**: 6 items
- Unhealthy Resources**: 6 items
- Unmonitored Resources**: 0 items

The resource list table includes columns for Resource name, Resource type, Subscription, Agent monitoring, Azure Defender, and Recommendations. Each row shows a status bar indicating the current state of these metrics.

Resource name	Resource type	Subscription	Agent monitoring	Azure Defender	Recommendations
dockeroniaasdemo	Container hosts	Contoso Hotels	On	Red	Red
dockervm-redhat	Container hosts	Contoso Hotels	On	Red	Red
acrtest	Container registries	Contoso Hotels	On	Green	Green
test-admission-controller	Kubernetes services	Contoso Hotels	On	Yellow	Yellow
yademo	Container registries	Contoso Hotels	On	Red	Red
m-kub-clust	Kubernetes services	Contoso Hotels	On	Yellow	Yellow

Azure Security Benchmark

Azure security benchmark is used here to organize these AKS security features.

- Network Security
- Identity Management
- Data Protection
- Logging & Monitoring
- Vulnerability

ref: <https://docs.microsoft.com/en-us/azure/security/benchmarks/overview>

Application Security

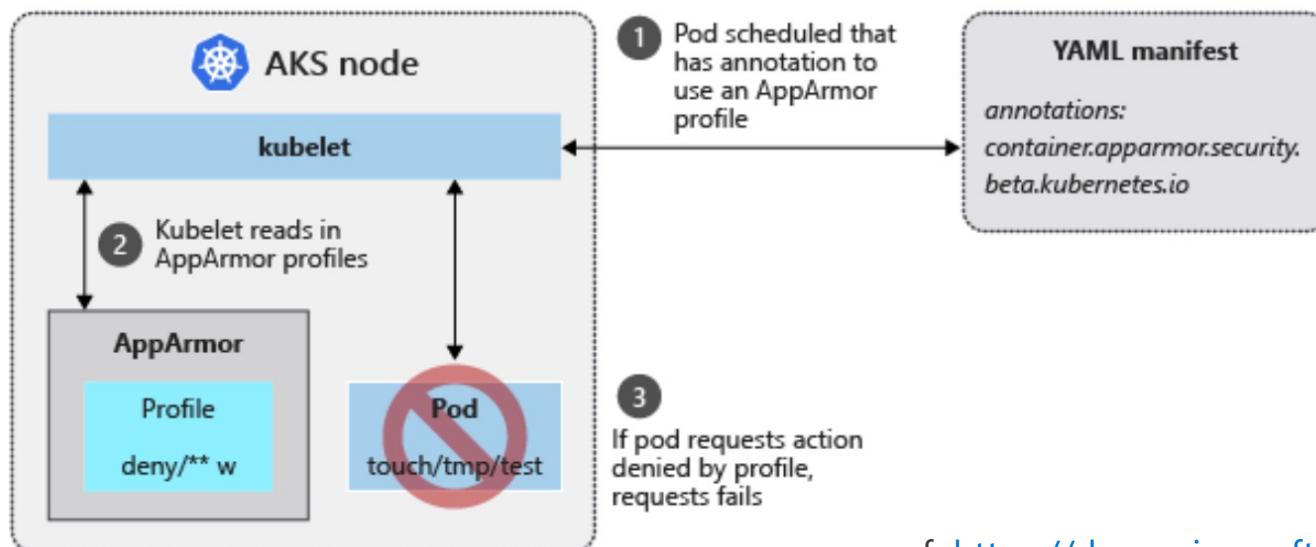
AppArmor in AKS

```
#include <tunables/global>
profile k8s-apparmor-example-deny-write flags=(attach_disconnected) {
    #include <abstractions/base>
    file,
    # Deny all file writes.
    deny /** w,
}
```

- Linux Kernel Security Module (like SELinux)
- *Enforcement or Complain mode*

annotations:

container.apparmor.security.beta.kubernetes.io/hello: localhost/k8s-apparmor-example-deny-write



Seccomp in AKS



SECCOMP
Linux Privilege Escalation

- Linux Kernel Security Module (like AppArmor)
- Works at the process level

annotations:

```
seccomp.security.alpha.kubernetes.io/pod: localhost/prevent-chmod
```

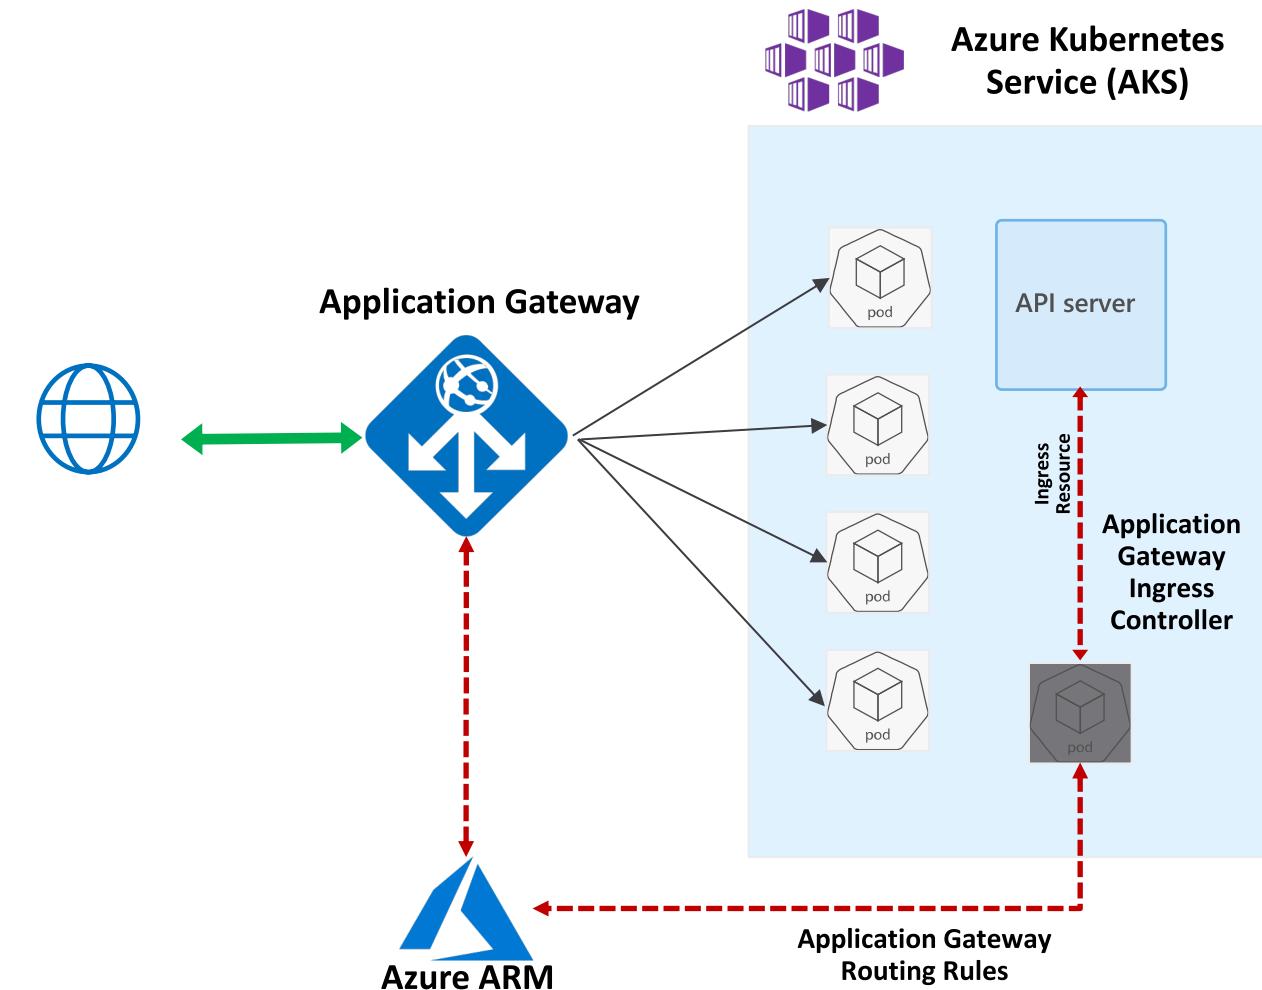
```
{  
    "defaultAction": "SCMP_ACT_ALLOW",  
    "syscalls": [  
        {  
            "name": "chmod",  
            "action": "SCMP_ACT_ERRNO"  
        },  
        {  
            "name": "fchmodat",  
            "action": "SCMP_ACT_ERRNO"  
        },  
        {  
            "name": "chmodat",  
            "action": "SCMP_ACT_ERRNO"  
        }  
    ]  
}
```

Network security

Application Gateway Ingress Controller

- Attach Application Gateways to AKS Clusters
- Load Balance from the Internet to pods (CNI and kubenet supported)
- Supports features of k8s ingress resource – TLS, multi-site and path-based routing
- Pod-AAD for ARM authentication

<https://github.com/Azure/application-gateway-kubernetes-ingress>





North/South Traffic in AKS with Azure Firewall

Cloud native stateful Firewall as a service

A first among public cloud providers

Target FQDN

*.azmk8s.io – http, https (Eg. *eastus.azmk8s.io)

k8s.gcr.io – http, https

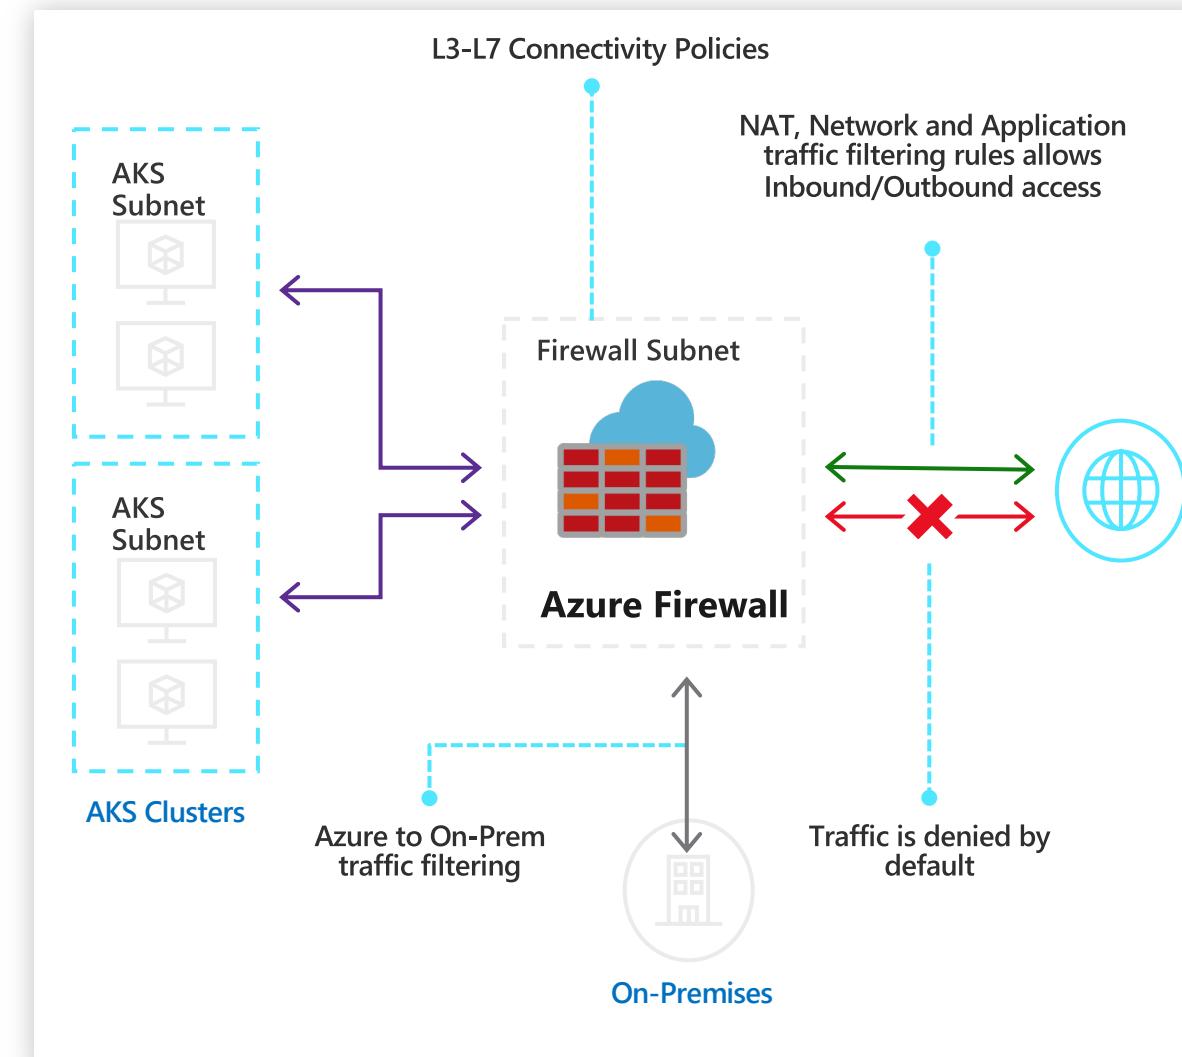
storage.googleapis.com - http, https

*auth.docker.io – http, https

*cloudflare.docker.io – http, https

*registry-1.docker.io – http, https

- NAT Rules
 - TCP Port 22 to AKS Destination Addresses for the Region

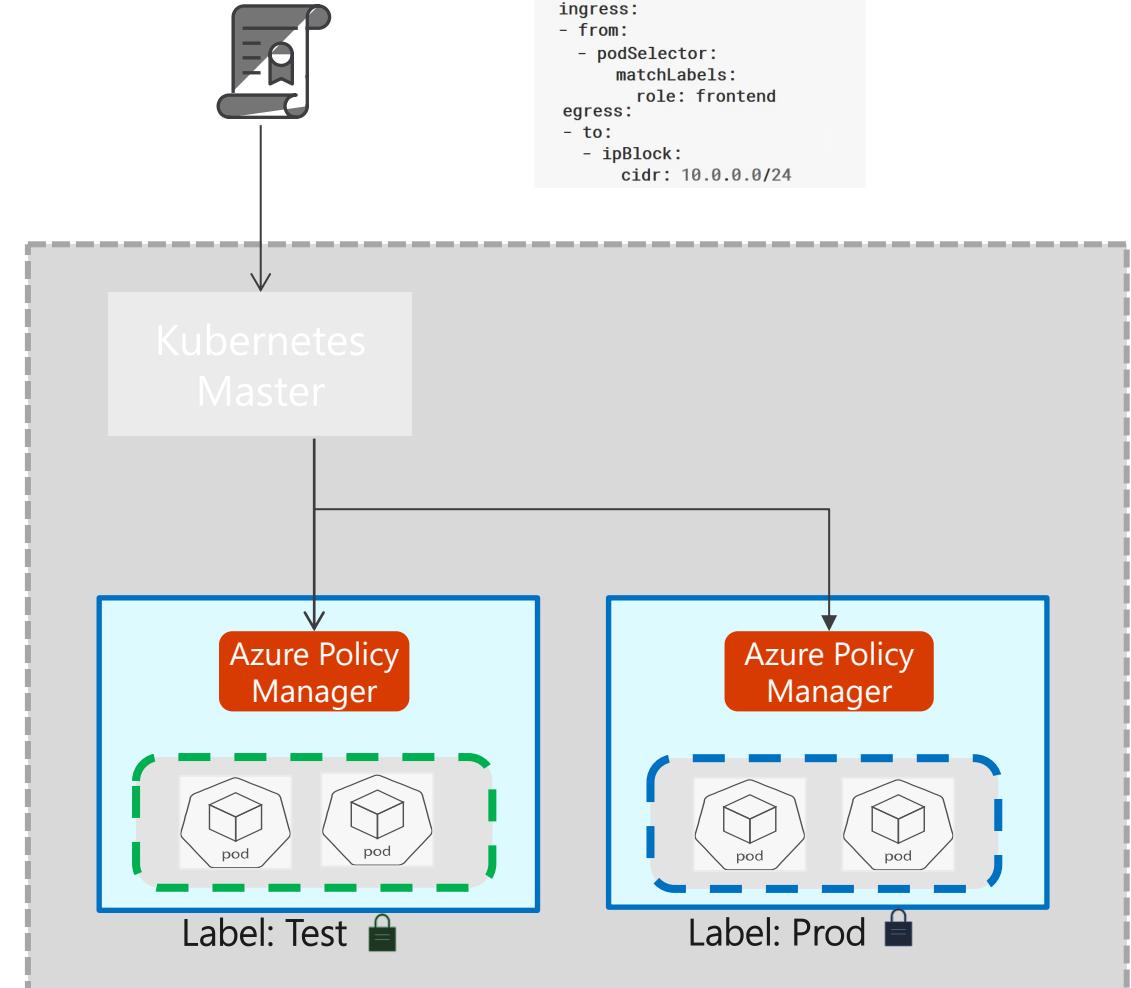


Network Policies

Provides micro-segmentation for containers
Label-based selection of Pods
Policy resource yaml file specifies Ingress and Egress policies
Works in conjunction with Azure CNI

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              role: frontend
  egress:
    - to:
        - ipBlock:
            cidr: 10.0.0.0/24
```

```
kubectl apply -f policy.yaml
```

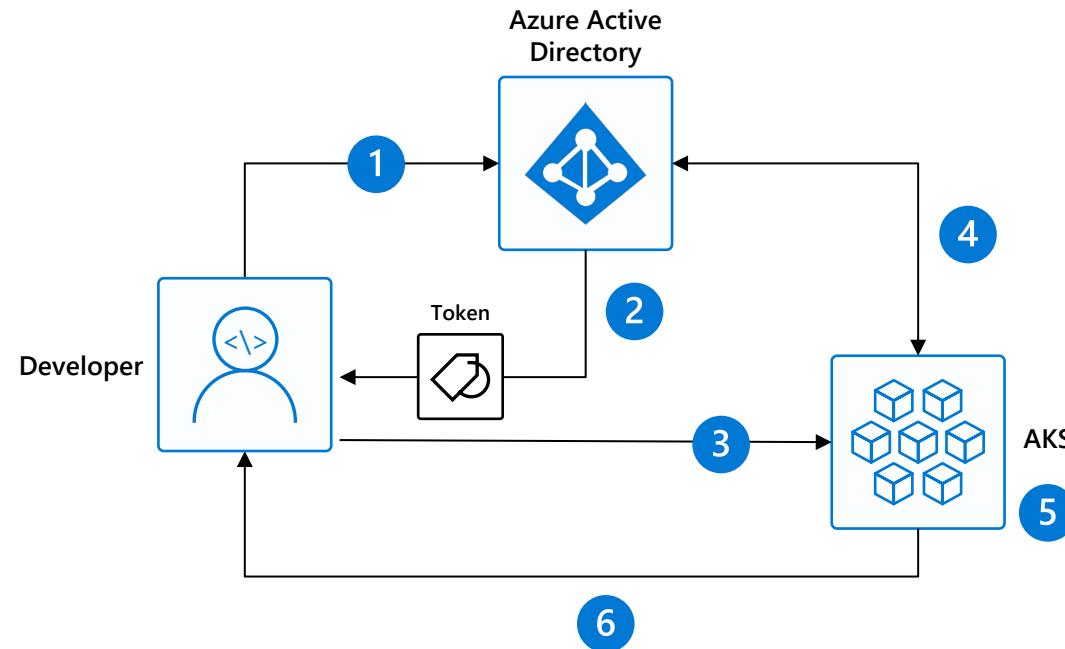


Kubernetes Cluster

Cluster access

Cluster Level – Kubernetes RBAC and Azure AD Authentication

1. Kubernetes Developer authenticates with AAD
2. The AAD token issuance endpoint issues the access token
3. Developer performs action w/ AAD token.
Eg. *kubectl create pod*
4. Kubernetes validates token with AAD and fetches the Developer's AAD Groups
Eg. Dev Team A, App Group B
5. Kubernetes RBAC and cluster policies are applied
6. Request is successful or not based on the previous validation



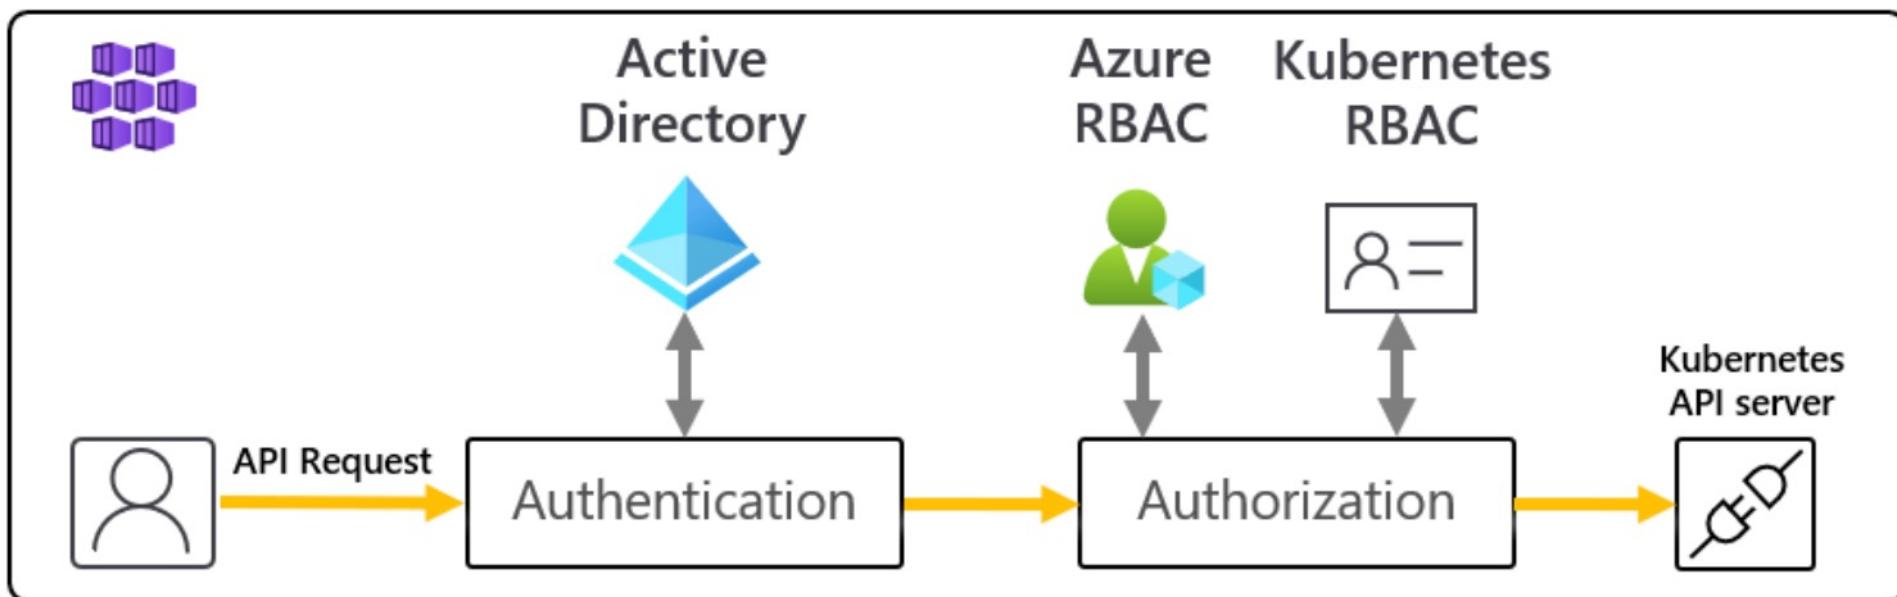
AAD with Azure RBAC Authorization

Usage: az aks create -g <rg> -n <name> --enable-aad --enable-azure-rbac

Roles:

- Azure Kubernetes Service RBAC Viewer,
- Azure Kubernetes Service RBAC Writer,
- Azure Kubernetes Service RBAC Admin,
- Azure Kubernetes Service RBAC Cluster Admin

```
az role assignment create --role "Azure Kubernetes Service RBAC Viewer" --assignee <aad-id> --scope <AKS_ID/namespaces/namespace-name>
```



AAD with Azure RBAC Authorization – custom roles

```
{  
  "Name": "AKS Deployment Reader and DS writer",  
  "Description": "Lets you view all deployments in cluster/namespace.",  
  "Actions": [],  
  "NotActions": [],  
  "DataActions": [  
    "Microsoft.ContainerService/managedClusters/apps/deployments/read",  
    "Microsoft.ContainerService/managedClusters/apps/statefulsets/write",  
  ],  
  "NotDataActions": [],  
  "assignableScopes": [  
    "/subscriptions/<YOUR SUBSCRIPTION ID>"  
  ]  
}
```

Upgrades

Node Vulnerability

Node update: Security updates are automatically applied to Linux nodes to protect customer's Azure Kubernetes Service (AKS) clusters. These updates include OS security fixes or kernel updates. AKS don't do reboot to complete the update.

Node image upgrade: AKS built VHD including OS and runtime update.
(including windows)

Usage: az aks upgrade -g <rg> -n <name> --node-image-only

Cluster Level – Nodes, Upgrade and Patches

Regular maintenance, security and cleanup tasks

- Maintain, update and upgrade hosts and kubernetes - Monthly ideal, 3 months minimum

◦ Security patches

- AKS automatically applies security patches to the nodes on a nightly schedule
- You're responsible to reboot as required, leveraging tools like Kured DaemonSet

<https://github.com/weaveworks/kured>

Upgrade to new versions

```
$ az aks upgrade --name myAKScluster \  
--resource-group myResourceGroup \  
--kubernetes-version 1.10.6
```

- **SSH Access**

- DenyEscalatingExec

- **Running benchmarks and tests to validate cluster setup**

- Kube-bench
- Aqua Hunter
- Others

CVE-2021-25741

kubernetes/kubernetes

#104980 **CVE-2021-25741: Symlink Exchange Can Allow...**



10 comments



cjcullen opened on September 13, 2021

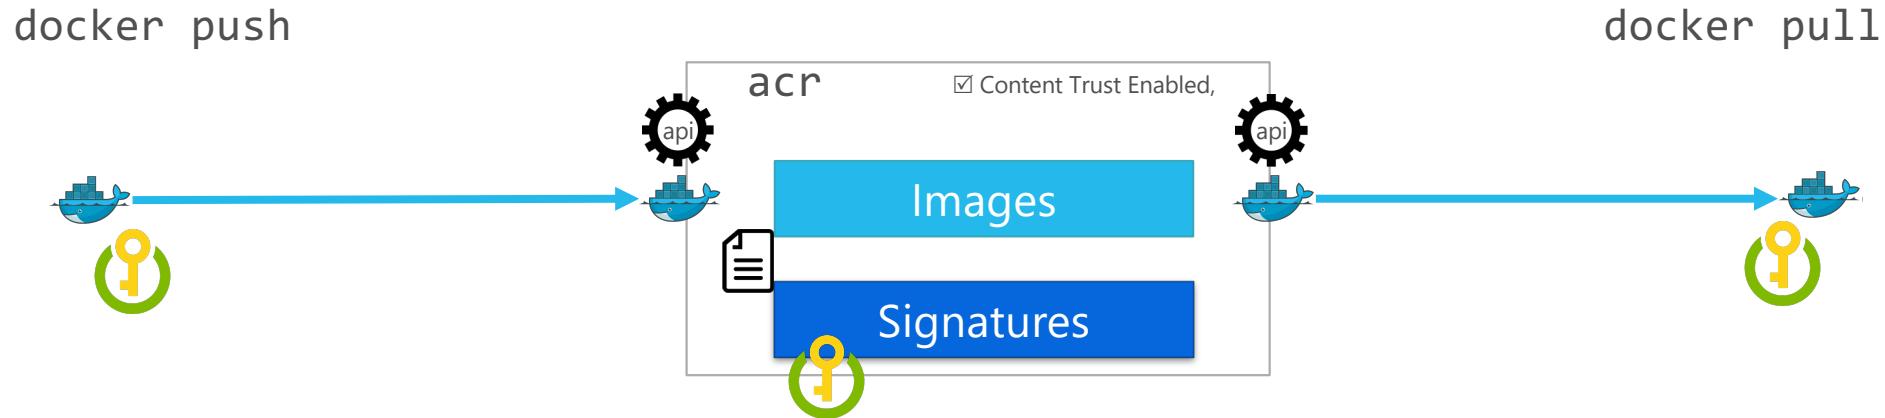


```
# az aks upgrade --resource-group rg-cluster --name cluster --node-image-only
```

[AKS/issues#2547](#)

Secure Supply Chain

ACR - Content Trust



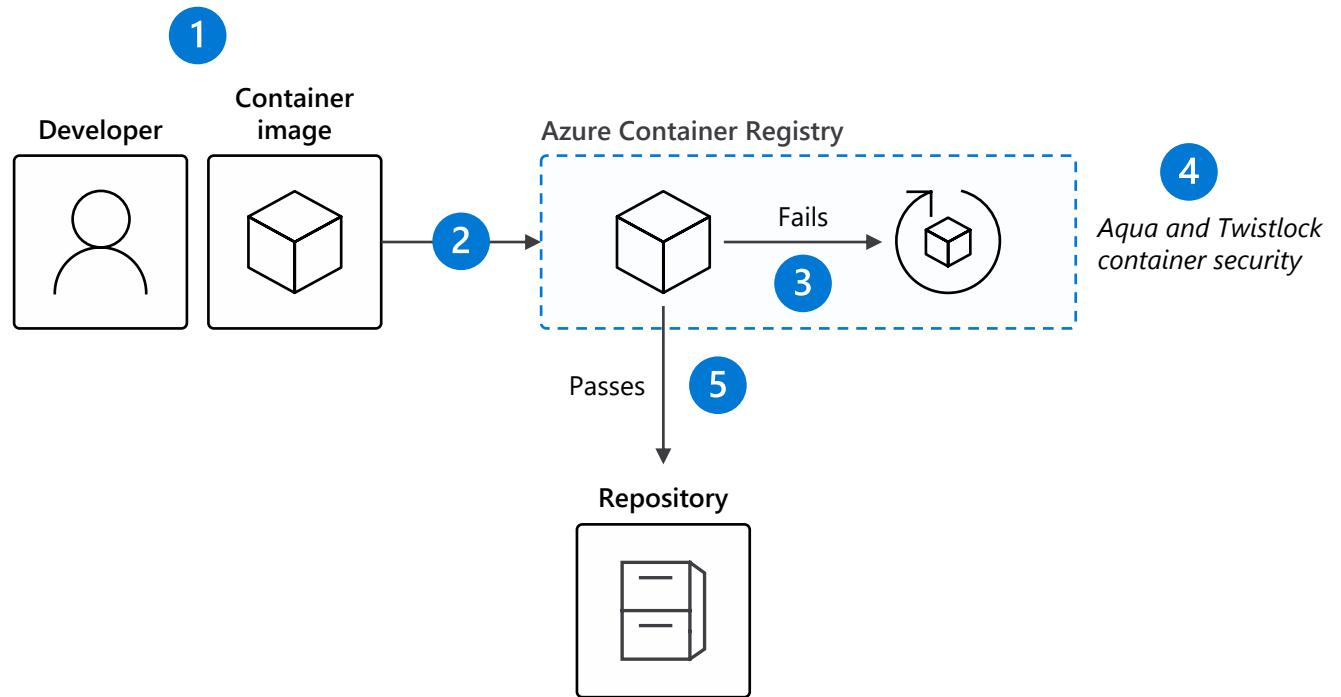
```
az role assignment create --scope <registry ID> --role AcrlImageSigner --assignee <user name>
```

To enforce on AKS: <https://github.com/sse-secure-systems/connaisseur>

Note: [notaryproject/notation](#) is the evolution of CT ([AKS/#688](#))

ACR – vulnerability scanning

1. Developer/CI system builds container image
2. Image pushed to Azure Container Registry
3. Azure Container Registry quarantines image until scanning passes
4. Azure Container Registry scans content leveraging Aqua, Twistlock
5. Azure Container Registry publishes the image to the repository



Azure Policy

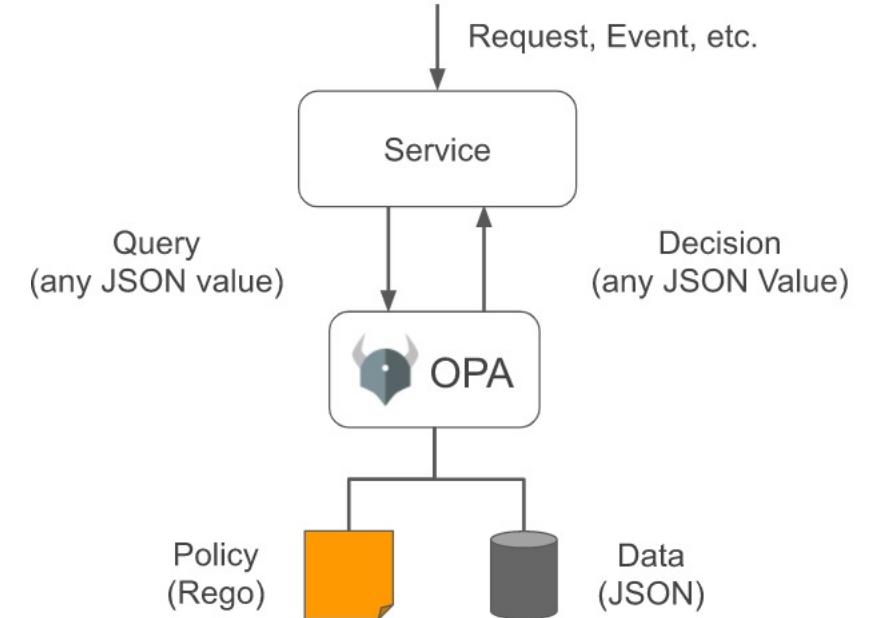
Azure Policy (open policy agent/gatekeeper)

The Azure Policy Add-on for AKS installs a managed instance of **Gatekeeper**, a validating/mutating admission controller.

Azure Policy for Kubernetes is built on the open-source Open Policy Agent.

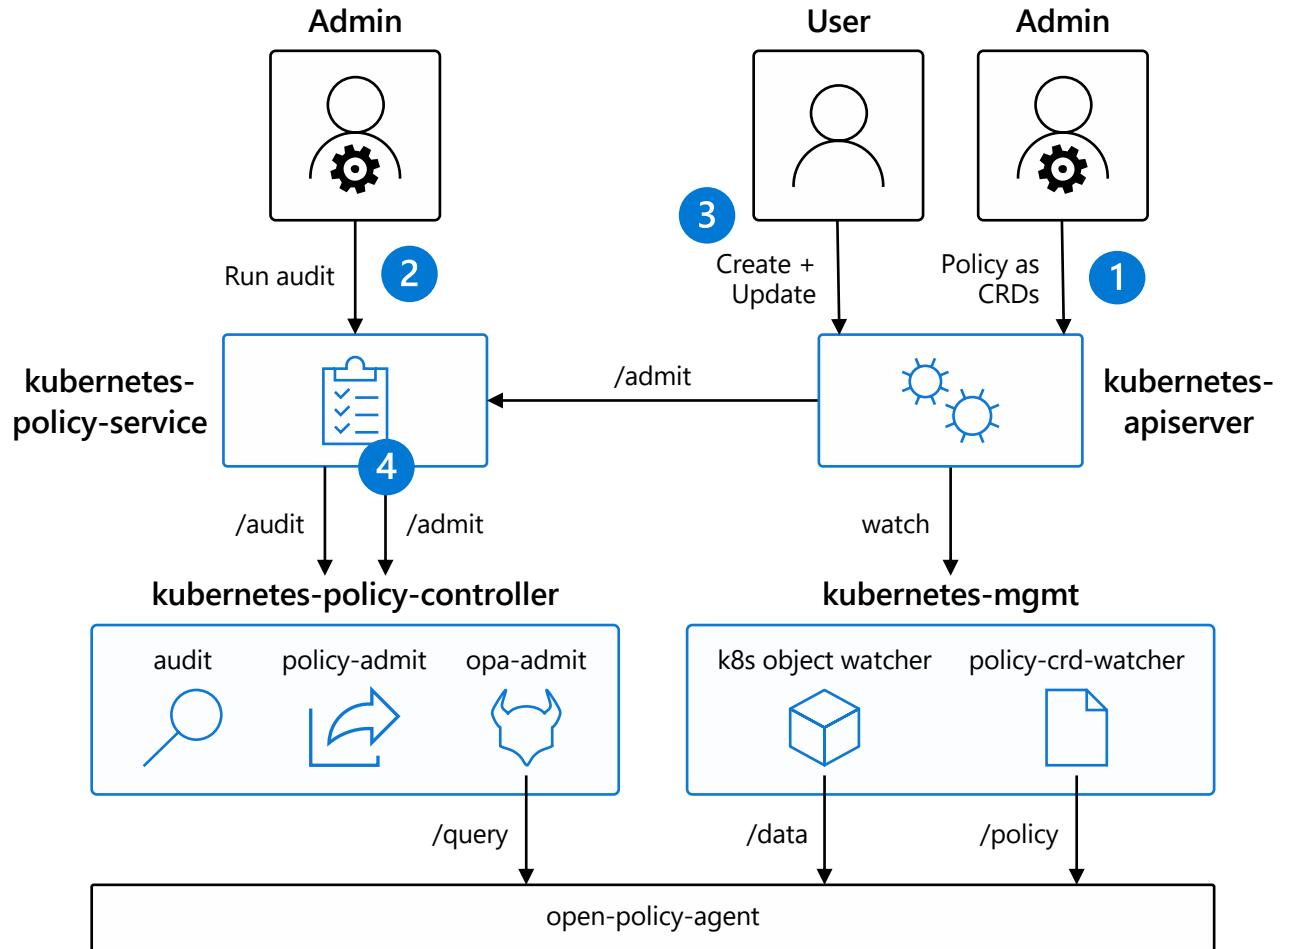
The add-on enacts the following functions:

- Checks with Azure Policy service for policy assignments to the cluster.
- Deploys policy definitions into the cluster as constraint template and constraint custom resources.
- Reports auditing and compliance details back to Azure Policy service.



Policy controller

1. Admin adds policy for the cluster
2. Admin audits compliance of the cluster using /audit endpoint
3. User uses standard Kubernetes API to operate the cluster and the create actions are guarded by policy
4. Kubernetes-policy-controller provides an admission controller webhook that performs evaluations by calling open-policy-agent (OPA) service



Azure Policy

```
apiVersion: templates.gatekeeper.sh/v1beta1
kind: ConstraintTemplate
metadata:
  name: k8sazurecontainernoprivilege
spec:
  crd:
    spec:
      names:
        kind: K8sAzureContainerNoPrivilege
        listKind: K8sAzureContainerNoPrivilegeList
        plural: k8sazurecontainernoprivilege
        singular: k8sazurecontainernoprivilege
  targets:
    - target: admission.k8s.gatekeeper.sh
      rego: |
        package k8sazurecontainernoprivilege

        violation[{"msg": msg, "details": {}}] {
          c := input_containers[_]
          c.securityContext.privileged
          msg := sprintf("Privileged container is not allowed: %v, securityContext: %v", [c.name, c.securityContext])
        }

        input_containers[c] {
          c := input.review.object.spec.containers[_]
        }

        input_containers[c] {
          c := input.review.object.spec.initContainers[_]
        }
```

```
apiVersion: constraints.gatekeeper.sh/v1beta1
kind: K8sAzureContainerNoPrivilege
metadata:
  name: container-no-privilege
spec:
  match:
    excludedNamespaces: {{ .Values.excludedNamespaces }}
    kinds:
      - apiGroups: []
        kinds: ["Pod"]
```

Azure Policy

 Now export your definitions and assignments to GitHub and manage them using actions! Click on 'Export definition' menu option. Learn more [here](#)

Name	↑↓	Definition loc...	↑↓	Policies	↑↓	Type	↑↓	Definition type	↑↓	Category
 Kubernetes cluster pod security restricted standards for Linux-based workloads				8		Built-in		Initiative		Kubernetes
 Kubernetes cluster pod security baseline standards for Linux-based workloads				5		Built-in		Initiative		Kubernetes
 [Preview]: Azure Policy Add-on for Kubernetes service (AKS) should be installed and enabled						Built-in		Policy		Kubernetes
 [Preview]: Deploy GitOps to Kubernetes cluster						Built-in		Policy		Kubernetes
 Both operating systems and data disks in Azure Kubernetes Service clusters should be encrypted						Built-in		Policy		Kubernetes
 Deploy Azure Policy Add-on to Azure Kubernetes Service clusters						Built-in		Policy		Kubernetes
 Kubernetes cluster pod hostPath volumes should only use allowed host paths						Built-in		Policy		Kubernetes
 Kubernetes cluster pods should only use allowed volume types						Built-in		Policy		Kubernetes
 Enforce HTTPS ingress in Kubernetes cluster						Built-in		Policy		Kubernetes
 Kubernetes clusters should not allow container privilege escalation						Built-in		Policy		Kubernetes

Secret Store CSI Driver

Azure Key Vault provider for Secret Store CSI Driver allows you to get secret contents stored in an Azure Key Vault instance and use the Secrets Store CSI driver interface to mount them into Kubernetes pods.

<https://github.com/Azure/secrets-store-csi-driver-provider-azure>

*Usage: az aks create –g <rg> -n <name>
install secret store csi driver into cluster
create your own secretProviderClass object
assign kubelet identity access to keyvault get access*

More info

<https://aka.ms/aks/releases>

<https://aka.ms/aks/roadmap>

<https://github.com/mspnp/aks-secure-baseline>

<https://www.the-aks-checklist.com/>

<https://azure.github.io/Aks-Construction/>



[Kubernetes Hardening Guide – NSA/CISA](#)